Reinforcement learning: An example

- Suppose you want to learn how to drive a car such that it follows a lane and minimises the number of collisions

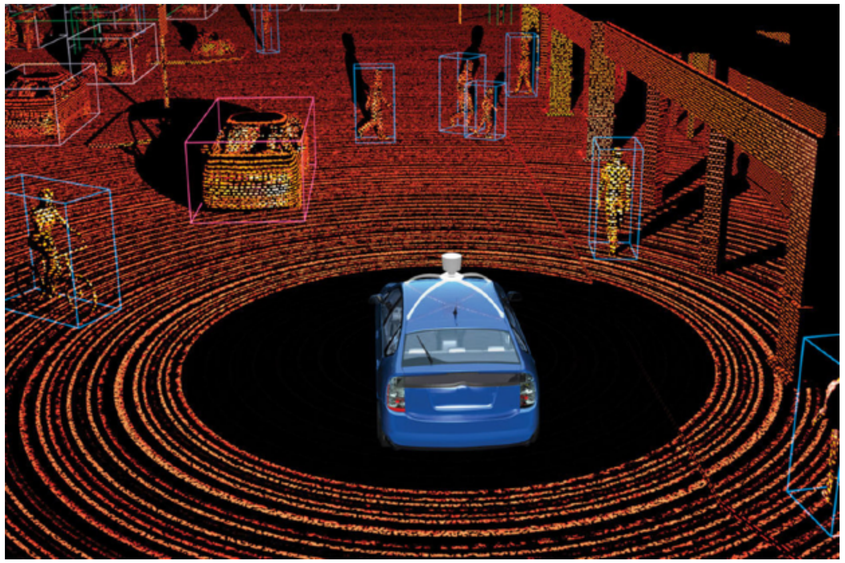

- You observe the state s of the surroundings: which can consist of very high dimensional data like LIDAR, images, etc.
- You take an action a = [steering angle, acceleration]
- Observe the next state $s' \sim \mathbb{P}(\cdot \,|\, s, a)$ and reward $= \mathbb{I}\{\text{on the lane}\} - \#\text{no of collisions}$
- Objective: Design a policy, which is a mapping from state to action e.g. a neural network, such that the reward over T rounds is maximized

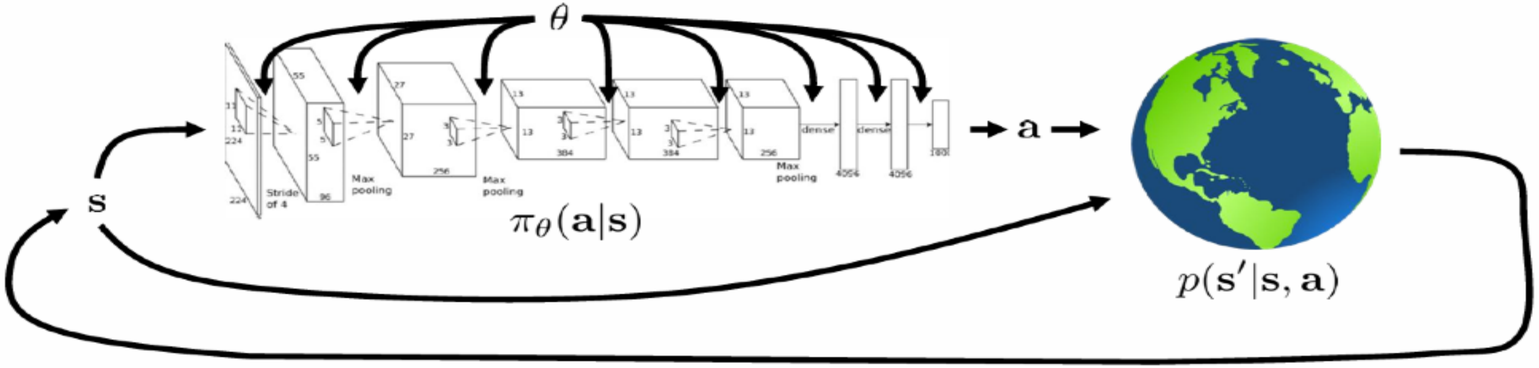

Reinforcement learning

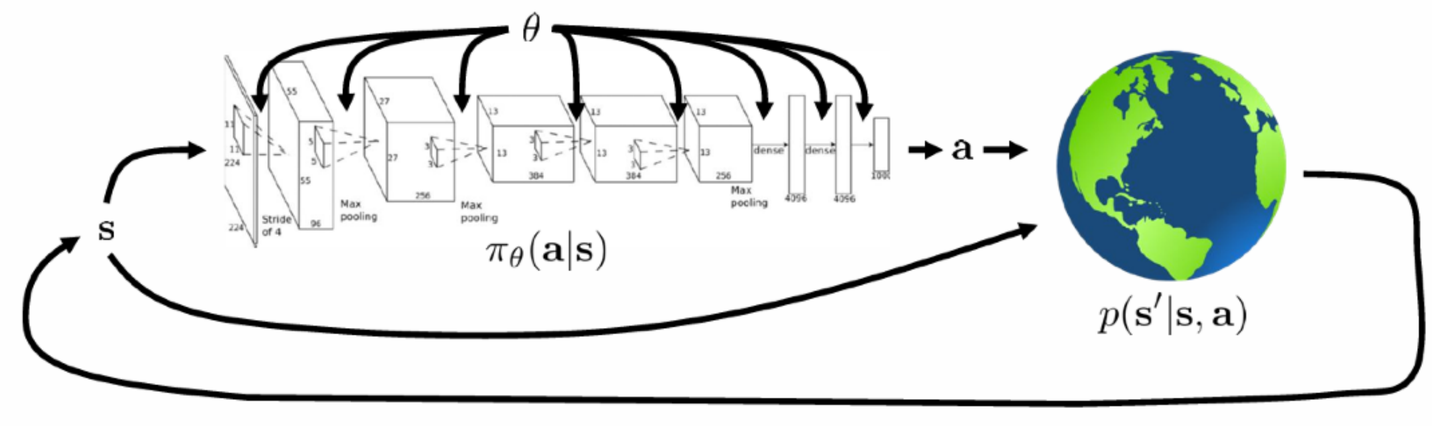

- Observe initial state s_0
- For every round $t = 1, \ldots, T$
  - Take action a_t
  - Observe next state $s_{t+1} \sim \mathbb{P}(\cdot \mid s_t, a_t)$ and reward r_t

Goal: Find policy π which is a mapping from state to action such that expected reward over T rounds is maximized

$$\sum_{t=1}^{T} \mathbb{E}_{a_t \sim \pi(\cdot|s_t), s_{t+1} \sim \mathbb{P}(\cdot|s_t, a_t)} [\gamma^{t-1} r_t(s_t, a_t) \mid s_0]$$